

AI in the Sciences and Engineering HS 2025: Lecture 10

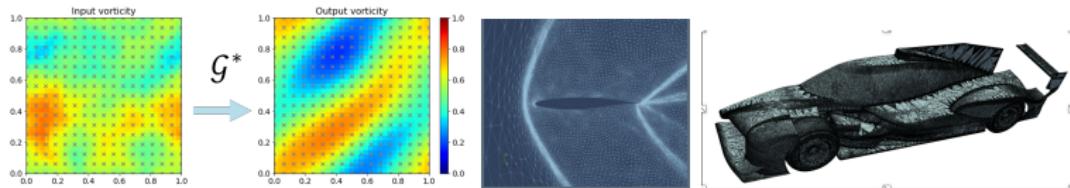
Siddhartha Mishra

Computational and Applied Mathematics Laboratory (CamLab)
Seminar for Applied Mathematics (SAM), D-MATH (and),
ETH AI Center (and) Swiss National AI Institute (SNAI) ,
ETH Zürich, Switzerland.

What we have learnt so far ?

- ▶ AIM: Learn PDEs using Deep Neural Networks
- ▶ Operator Learning: Learn the PDE Solution Operator from data.
- ▶ Examples: FNO, CNO, VIT, scOT etc.
- ▶ Very successful on PDEs on 2D Cartesian Domains !!
- ▶ Readily extended to time-dependent PDEs with all2all training.
- ▶ What are the caveats ?

Caveat II: PDEs on Arbitrary Domains

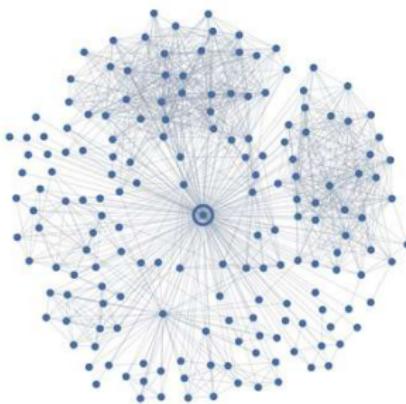


- ▶ Discussion so far has only focussed on **Cartesian Domains**
- ▶ Discretized with **Uniform Grids**.
- ▶ Most Real world PDEs are on **Arbitrary Domains**
- ▶ Discretized with **Unstructured Grids** or **Point Clouds**
- ▶ Need to handle such Data !!

Summary of Available Methods

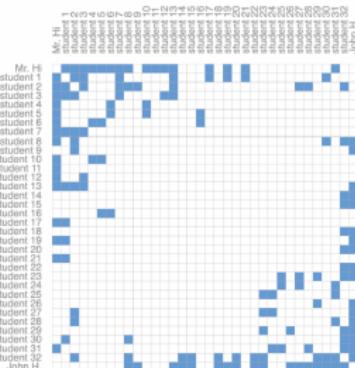
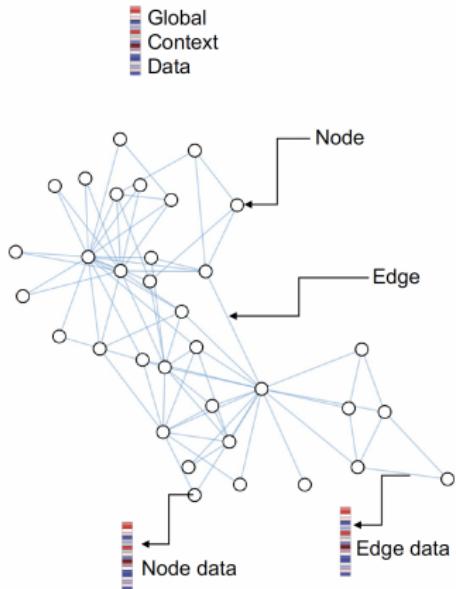
- ▶ Posing problem on Cartesian Domain through Masking.
- ▶ Direct Spectral Evaluations for FNO
 - ▶ DSE proposed in Lingsch, SM et. al., 2024
- ▶ Graph Neural Networks

Learning Graph Structured data



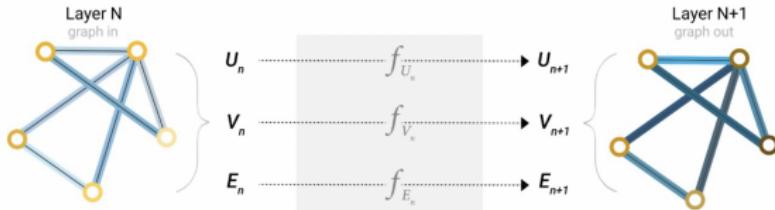
- ▶ Input data on **Graphs** is ubiquitous:
 - ▶ Social networks, Recommender and Transport systems.
 - ▶ Chemistry and Biology (Molecules)
 - ▶ Mesh-based numerical simulations in Scientific computing

Basics on Graphs



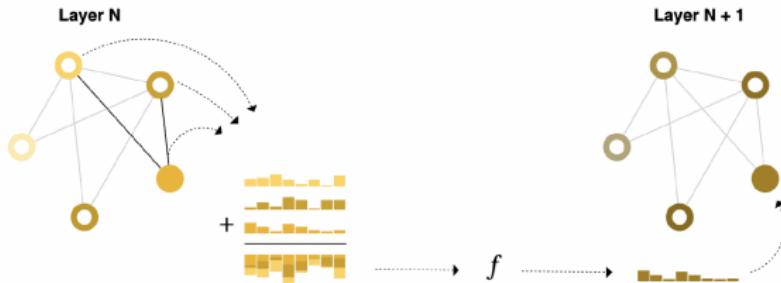
Adjacency matrix

What is a Graph Neural Network ?



- ▶ GNN is a composition of layers.
- ▶ Each layer updates Node + Edge features

Message Passing is the basis of all GNNs

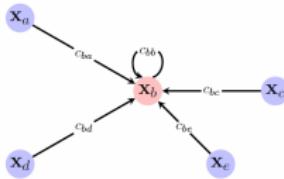


- ▶ Generic form of Message Passing:

$$h_i := f \left(v_i, \bigoplus_{j \in \mathcal{N}_i} \Psi(v_i, v_j) \right)$$

- ▶ f, Ψ are MLPs.
- ▶ \bigoplus is an aggregation function

Ex I: Graph Convolutional Network (GCN)



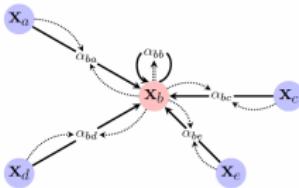
- ▶ Generic form of **GCN**:

$$h_i := f \left(v_i, \bigoplus_{j \in \mathcal{N}_i} c_{ij} \Psi(v_j) \right)$$

- ▶ Specific example:

$$h_i := g \left(\sum_{j \in \mathcal{N}_i \cup \{v_i\}} \frac{1}{\sqrt{\tilde{d}_i \tilde{d}_j}} W v_j \right), \quad \tilde{d}_i = 1 + \sum_j A_{ij}$$

Ex II: Graph Attention Network (GAT)



- ▶ Generic form of GAT:

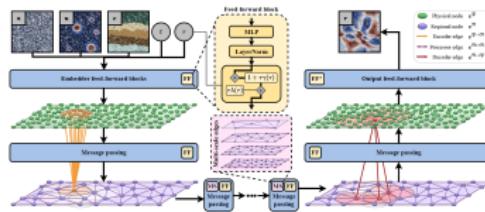
$$h_i := f \left(v_i, \bigoplus_{j \in \mathcal{N}_i} \alpha(v_i, v_j) \psi(v_j) \right)$$

- ▶ Specific example:

$$h_i := g \left(\sum_{j \in \mathcal{N}_i \cup \{v_i\}} \alpha(v_i, v_j) W v_j \right),$$

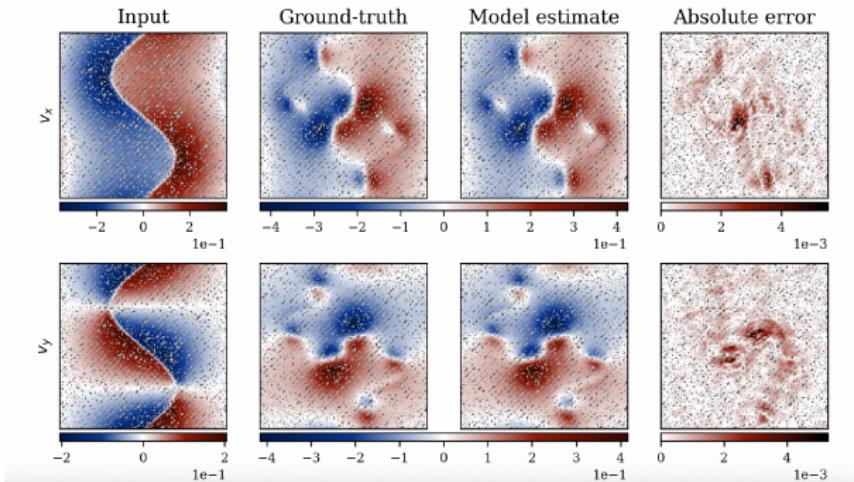
- ▶ α is a softmax based weight.

GNN for Operator Learning of PDEs ?



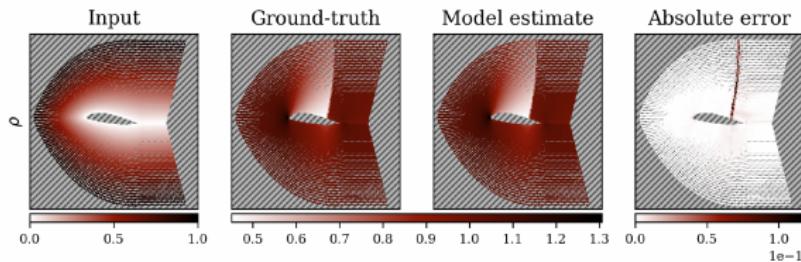
- ▶ **RIGNO** of [Mousavi, SM, et. al., 2025](#)
- ▶ Based on general [MPNNs](#)
- ▶ Multiple modifications to ensure:
 - ▶ Multiscale Information processing.
 - ▶ Temporal Continuity
 - ▶ Resolution Invariance

NS-SVS



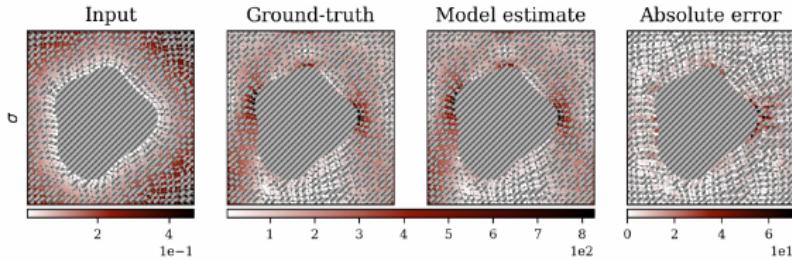
Model	FNO-DSE	GeoFNO	GINO	RIGNO
Errors (in %)	26.0	9.75	1.19	0.56

Flow Past Airfoils

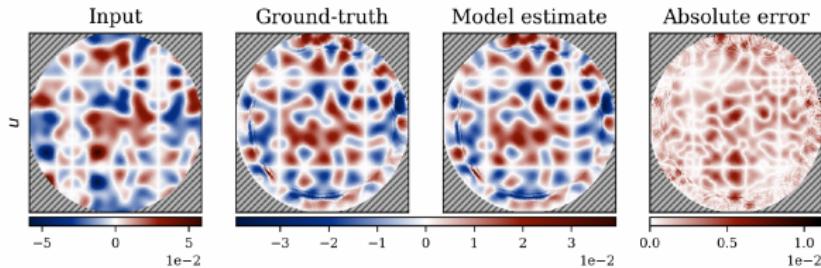


Model	FNO-DSE	GeoFNO	GINO	RIGNO
Errors (in %)	1.99	4.48	2.00	1.09

Plasticity



Waves in a Circle



Model	FNO-DSE	GeoFNO	GINO	RIGNO
Errors (in %)	5.52	13.1	5.82	5.35

RIGNO is very accurate

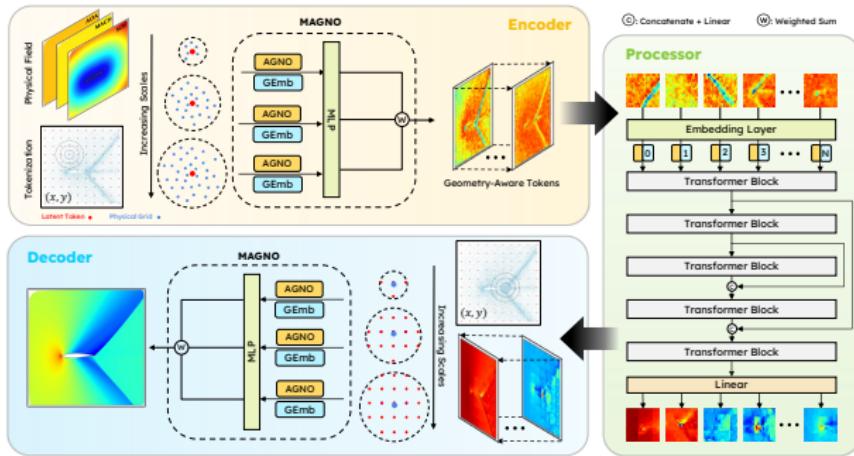
Dataset <i>(unstructured)</i>	Median relative L^1 error [%]				
	RIGNO-18	RIGNO-12	GeoFNO	FNO DSE	GINO
Heat-L-Sines	0.04	0.05	0.15	0.53	0.19
Wave-C-Sines	5.35	6.25	13.1	5.52	5.82
NS-Gauss	2.29	3.80	41.1	38.4	13.1
NS-PwC	1.58	2.03	26.0	56.7	5.85
NS-SL	1.28	1.91	24.3	29.6	4.48
NS-SVS	0.56	0.73	9.75	26.0	1.19
CE-Gauss	6.90	7.44	42.1	30.8	25.1
CE-RP	3.98	4.92	18.4	27.7	12.3
ACE	0.01	0.01	1.09	1.29	3.33
Wave-Layer	6.77	9.01	11.1	28.3	19.2
AF	1.00	1.09	4.48	1.99	2.00
Elasticity	4.31	4.63	5.53	4.81	4.38
<i>(uniform grid)</i>	RIGNO-18	RIGNO-12	CNO	scOT	FNO
	NS-Gauss	2.74	3.78	10.9	2.92
NS-PwC	1.12	1.82	5.03	7.11	11.24
NS-SL	1.13	1.82	2.12	2.49	2.08
NS-SVS	0.56	0.75	0.70	0.99	6.21
CE-Gauss	5.47	7.56	22.0	9.44	28.69
CE-RP	3.49	4.43	18.4	9.74	31.19
ACE	0.01	0.01	0.28	0.21	0.60
Wave-Layer	6.75	8.97	8.28	13.44	28.01

What is the Issue then ?



- ▶ But **not Efficient** (more than 2d of training) !!
- ▶ Due to **Repeated Sparse Memory Access**

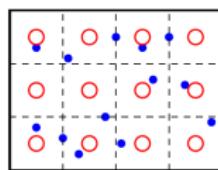
Solution ?: Use Graphs + Transformers



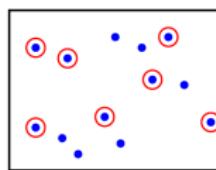
- ▶ **Geometry Aware Operator Transformer:** Wen, SM et. al, 2025

On GAOT

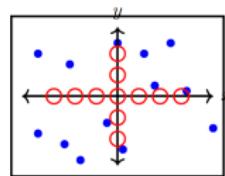
- ▶ Based on the **Encode-Process-Decode** Strategy.
- ▶ **Encoder:** Input Point Cloud \mapsto **Latent Grid**
- ▶ Different choices of Latent Grid:



(a) Strategy I



(b) Strategy II



(c) Strategy III

GAOT Encoder

- Given input field a , Encoder outputs w_e at any Latent point y

$$w_e(y) = \sum_{y=1}^{n_y} \alpha_k K(y, x_k, a(x_k)) \varphi(a(x_k)), \quad |y - x_k| \leq r$$

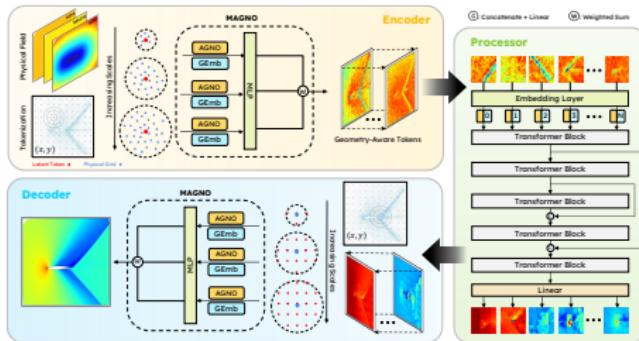
- MLPs: K, φ
- Quadrature weights α chosen by Attention:

$$\alpha_k = \frac{\exp(e_k)}{\sum_{k'=1}^{n_y} \exp(e_{k'})}, \quad e_k = \frac{\langle W_q y, W_k x_k \rangle}{\sqrt{d}}$$

- We use Multiscale Aggregation over m scales:

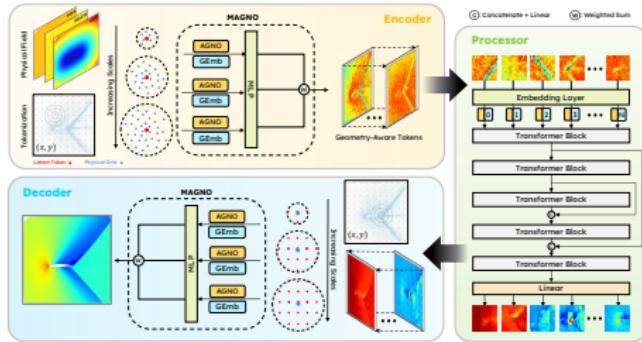
$$w_e^m(y) = \sum_{y=1}^{n_y^m} \alpha_k^m K^m(y, x_k, a(x_k)) \varphi^m(a(x_k)), \quad |y - x_k| \leq r_m$$

What about the Geometry ?



- ▶ Geometric Information is conveyed through:
 - ▶ Point coordinates
 - ▶ Signed Distance functions (SDFs)
 - ▶ Local Statistical Embeddings $g_m(y)$ based on,
 - ▶ Number of Neighbors,
 - ▶ Average Distance + its Variance
 - ▶ Local Anisotropy through PCs of the distance covariance matrix

MAGNO Encoder

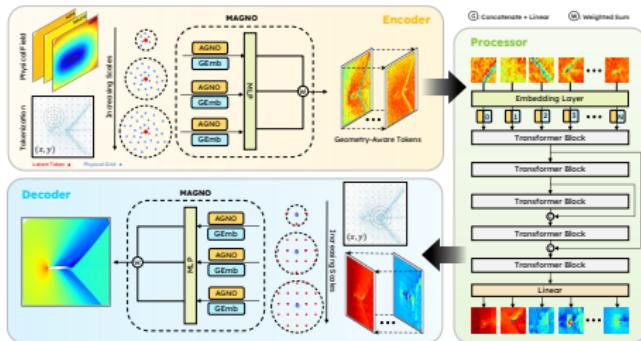


- ▶ Final form of the Encoder:

$$w_e(y) = \sum_{m=1}^M \beta_k^m \hat{w}^m(y), \quad \beta_m(y) = \frac{\exp(\Psi_m(y))}{\sum_{m'=1}^M \exp(\Psi_{m'}(y))}$$

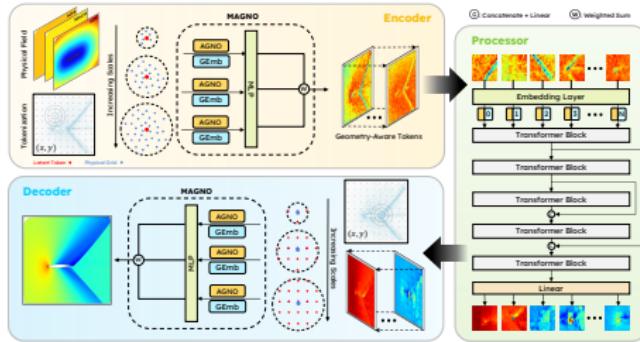
- ▶ Provide a set of **Geometry Aware Tokens**

GOAT Processor



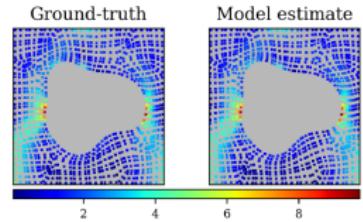
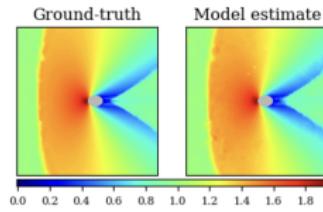
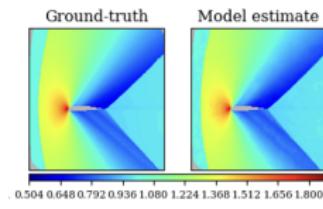
- ▶ Use of ViT or SWIN for Cartesian Latent Grids.
- ▶ Allows for Patching
- ▶ Standard seq2seq Transformers for Unstructured Latent Grids.

GOAT Decoder



- ▶ MAGNO Decoder
- ▶ Allows evaluation of output at any query point $y \in D$
- ▶ ensures Neural Field Property.
- ▶ GINO:
 - ▶ Single scale GNO encoder + decoder
 - ▶ FNO processor.

GAOT is very flexible wrt Domain Geometry

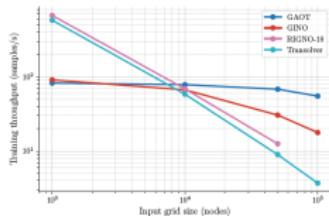


GAOT is vey accurate and robust

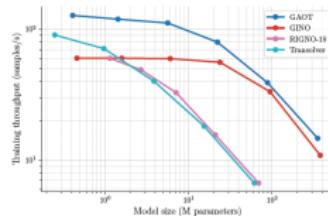
Dataset		Median relative L^1 error [%]				
Time-Independent	GAOT	RIGNO-18	Transolver	GNOT	UPT	GINO
Poisson-C-Sines	3.10	6.83	77.3	100	100	20.0
Poisson-Gauss	0.83	2.26	2.02	88.9	48.4	7.57
Elasticity	1.34	4.31	4.92	10.4	12.6	4.38
NACA0012	6.81	5.30	8.69	6.89	16.1	9.01
NACA2412	6.66	6.72	8.51	8.82	17.9	9.39
RAE2822	6.61	5.06	4.82	7.15	16.1	8.61
Bluff-Body	2.25	5.76	1.78	44.2	5.81	3.49
Time-Dependent	GAOT	RIGNO-18	GeoFNO	FNO DSE	UPT	GINO
NS-Gauss	2.91	2.29	41.1	38.4	92.5	13.1
NS-PwC	1.50	1.58	26.0	56.7	100	5.85
NS-SL	1.21	1.28	13.7	22.6	51.5	4.48
NS-SVS	0.46	0.56	9.75	26.0	4.2	1.19
CE-Gauss	6.40	6.90	42.1	30.8	64.2	25.1
CE-RP	5.97	3.98	18.4	27.7	26.8	12.3
Wave-Layer	5.78	6.77	11.1	28.3	19.6	19.2
Wave-C-Sines	4.65	5.35	13.1	5.52	12.7	5.82

GAOT is also very Efficient

- ▶ In terms of
- ▶ Training Throughput
- ▶ Inference Latency
- ▶ wrt Model Size + Input Size.
- ▶ This is due to Efficient Implementation



(a) Grid vs. Throughput



(b) Model vs. Throughput

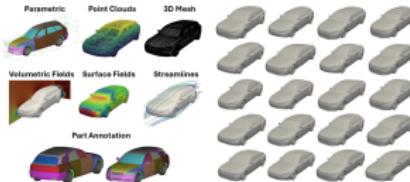
GAOT combines best of both Worlds



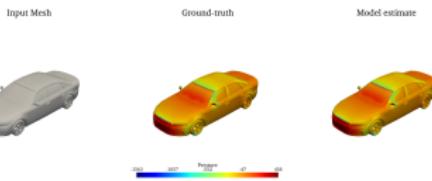
- We can scale it to 3d Industrial Scale datasets.

The DrivaerNet++ Benchmark

- ▶ Flow past Cars Dataset (8K Car Shapes with 0.5M nodes)



- ▶ GAOT: SOTA for Surface Pressure, Shear Stress

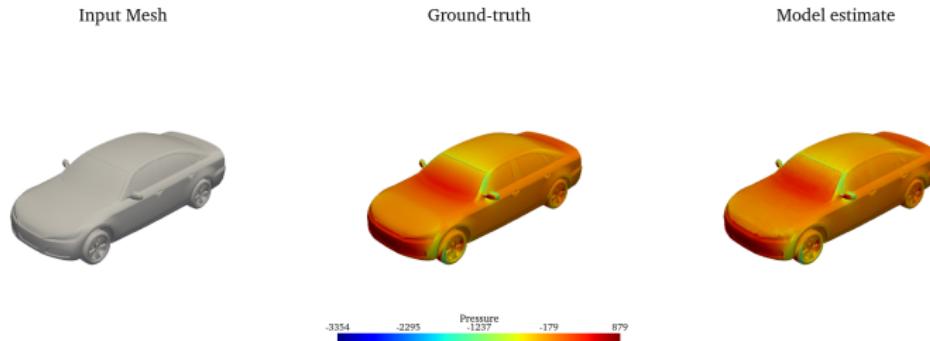


Model	GAOT	FigConvNet	TripNet	RegDGCNN
L^1 Pressure Err.	0.110	0.122	0.125	0.161
L^1 Shear Err.	0.156	0.222	0.215	0.364

- CFD: 375 Node hours vs. GAOT: 0.36 seconds !!!

The DrivaerML Benchmark

- ▶ HR-LES simulations of flow past 500 cars.
- ▶ More accurate than RANS for Drivearnet++.
- ▶ 9 M surface nodes handled accurately by GAOT !!



	Model	GAOT	GINO
L^1 Pressure Err.	0.123	0.152	
L^1 Shear Err.	0.216	0.273	

- CFD: 61K Node hours vs. GAOT: 14 seconds !!!

Flow Past an entire Aircraft

- ▶ AIAA's **NASA CRM** Benchmark.
- ▶ **GAOT** predicts Surface Pressure+Skin Friction accurately !!

